

TECHNOVA SOLUTIONS

Official Service Tier & Cloud Architecture Pricing Guide

TechNova Solutions provides highly scalable cloud infrastructure, optimized data pipelines, and intelligent agentic orchestration platforms. This Pricing Guide outlines our standard service tiers built to handle diverse computational and organizational requirements.

1. Starter Plan

Engineered for individual developers, innovators, and early prototyping environments seeking a robust technical foundation.

Pricing: \$29/mo (Monthly) | \$24/mo billed as \$288/yr (Yearly)

- **Storage Capacity:** 50 GB Secure SSD Cloud Storage
- **AI Compute Requests:** Up to **10,000** model inferences per billing cycle
- **Team Provisions:** 1 Primary Administrator seat
- **Support Ecosystem:** Standard Electronic Ticket support Response window ≤ 24 hours
- **API Infrastructure:** Sandbox Rate-limited access (Max 60 requests/min)
- **Analytics Layer:** Core usage visualization metrics dashboard
- **Security Architecture:** AES-256 baseline encryption at rest and TLS 1.3 in transit

2. Professional Plan

Tailored for expanding teams, agile startups, and professional content operations needing low-latency runtimes and expanded quotas.

Pricing: \$99/mo (Monthly) | \$79/mo billed as \$948/yr (Yearly)

- **Storage Capacity:** 500 GB High-Throughput NVMe Storage
- **AI Compute Requests:** Up to **100,000** model inferences per billing cycle
- **Team Provisions:** Up to 5 Collaborative Seats
- **Support Ecosystem:** Priority Email Desk support Response window ≤ 4 hours
- **API Infrastructure:** Production-grade access with dedicated quotas (Max 300 requests/min)
- **Analytics Layer:** Advanced custom query dashboards and scheduled reports export
- **Security Architecture:** Multi-Factor Authentication (MFA) enforcement and role-based access control (RBAC)

3. Business Plan

Designed for mid-market commercial operations demanding persistent database nodes, optimized pipeline execution, and broader collaboration limits.

Pricing: \$299/mo (Monthly) | \$239/mo billed as \$2,868/yr (Yearly)

- **Storage Capacity:** 2 TB Dedicated Cloud Storage Infrastructure
- **AI Compute Requests:** Up to **500,000** model inferences per billing cycle
- **Team Provisions:** Up to 20 Collaborative Seats

- **Support Ecosystem:** 24/7 Priority Support (Live Chat & Dedicated Help Desk Access)
- **API Infrastructure:** Fully unthrottled endpoints with dedicated custom routing pipelines
- **Analytics Layer:** Real-time streaming operational telemetry and programmatic webhook exports
- **Security Architecture:** Single Sign-On (SSO / SAML 2.0) integration and SOC 2 Type II audit log streaming

4. Enterprise Plan

Crafted for multi-tenant corporations, institutional scale operations, and heavy machine learning workflows requiring custom limits and dedicated resources.

Pricing: Custom Quotation Billed Annually (Contact Sales Architecture Team)

- **Storage Capacity:** Completely Uncapped / Custom allocated Distributed Cloud Architecture
- **AI Compute Requests:** Unlimited high-speed fine-tuned inferences via dedicated dedicated model nodes
- **Team Provisions:** Unlimited Seat allocations across multiple parent and child sub-organizations
- **Support Ecosystem:** Dedicated Systems Engineering Support Matrix with a *15 ext{-min}* critical SLA
- **API Infrastructure:** Custom, fully isolated API mesh networks tailored to specific high-velocity processing
- **Analytics Layer:** Direct data-lake syncing pipelines, custom ML metrics tracking, and raw event data exporting
- **Security Architecture:** Advanced IP Whitelisting, Private Virtual Clouds (VPC), Hardware Security Module (HSM) level keys, and full HIPAA / GDPR baseline configurations

5. Comprehensive Matrix Comparison Table

Feature / Metric	Starter	Professional	Business	Enterprise
Monthly Cost	\$29 / mo	\$99 / mo	\$299 / mo	Custom Quote Only
Yearly Commuted Cost	\$24 / mo (\$288 total)	\$79 / mo (\$948 total)	\$239 / mo (\$2,868 total)	Custom Bespoke Billing
Storage Volume	50 GB Secure SSD	500 GB High-Speed NVMe	2 TB Dedicated Volume	Uncapped / Tailored
AI Requests Limit	10,000 / mo	100,000 / mo	500,000 / mo	Completely Unlimited
Max Seats (Team)	1 User Node	5 Collaborative Seats	20 Collaborative Seats	Unlimited Allocation
Support Level Matrix	Standard Ticket (≤ 24h)	Priority Email (≤ 4h)	24/7 Live Desk Support	Dedicated Account Engineer Matrix
API Access Provision	Sandbox Rate-Limited	Production (300 req/min)	Unthrottled High-Volume	Custom Isolated API Mesh Networks
Analytics Capability	Core Visualization	Advanced / Custom Queries	Real-time Telemetry Streams	Data-Lake Infrastructure Sync
Security Framework	AES-256 & TLS 1.3	MFA Enforcement & RBAC	SSO / SAML 2.0 & Audit Logs	Isolated VPC, HSM Keys, Custom Compliance